



# Complete Testing & Testcontainers Guide

---

The ONLY guide you need for testing with Spring Boot, Testcontainers, and PostgreSQL

**Last Updated:** December 31, 2025

---



## Table of Contents

### Part 1: Quick Start

1. [TL;DR - Run Tests in 3 Steps](#)
2. [What's in Your Test Setup](#)

### Part 2: Setup & Configuration

3. [Prerequisites](#)
4. [Global Configuration Files](#)
5. [Test Configuration](#)

### Part 3: Writing Tests

6. [Basic Test Structure](#)
7. [Data Injection Tests](#)
8. [Test Execution Order](#)

### Part 4: Database Management

9. [Connecting to Test Database](#)
10. [Container Lifecycle](#)
11. [Stopping Containers](#)

### Part 5: Understanding Your Setup

12. [Spring vs Prisma](#)
13. [JPA Foreign Keys](#)
14. [SQLite vs PostgreSQL in Tests](#)
15. [Tables Dropping After Tests](#)

### Part 6: Troubleshooting

16. [Common Issues](#)
  17. [FAQ](#)
- 

## Part 1: Quick Start

---

### TL;DR - Run Tests in 3 Steps

```
# 1. Start Docker
open -a Docker # macOS

# 2. Run all tests
./gradlew test

# 3. Run test suite with data injection
./gradlew test --tests IntegrationTestSuite
```

**That's it!** PostgreSQL container starts automatically, schema generated from JPA entities, tests run.

---

## What's in Your Test Setup

### Your Test Stack

- **Spring Boot 3.2** with Kotlin
- **Testcontainers** for PostgreSQL
- **JPA/Hibernate** for schema generation
- **16 JPA Entities** with relationships
- **Container reuse** for fast tests
- **Test execution order** control

### Production vs Tests

Aspect	Production	Tests
<b>Database</b>	SQLite	PostgreSQL (Testcontainers)
<b>Schema</b>	Managed by you	Auto-generated from entities
<b>Port</b>	Fixed	Dynamic (e.g., 52106)
<b>Data</b>	Persistent	Ephemeral (per test run)

---

## Part 2: Setup & Configuration

---

### Prerequisites

#### Required

1. **Docker Desktop** - Must be running

```
open -a Docker # Start Docker
docker ps      # Verify it's running
```

2. **JDK 17** - Already configured

### 3. Gradle - Included (./gradlew)

Optional

- Node.js/npm (not needed for tests, only for Prisma in production)
- 

## Global Configuration Files

### 1. ~/.testcontainers.properties

**Location:** Your home directory (`~/.testcontainers.properties`)

**Content:**

```
testcontainers.reuse.enable=true
```

**What it does:**

- Enables container reuse across test runs
- Makes subsequent test runs much faster (seconds vs minutes)
- Container stays alive until Docker restart

**Create it:**

```
echo "testcontainers.reuse.enable=true" > ~/.testcontainers.properties
```

**Or use the script:**

```
./create-testcontainers-config.sh
```

---

### 2. src/test/resources/application-test.yml

**Content:**

```
spring:  
  jpa:  
    hibernate:  
      ddl-auto: create # Creates tables, keeps after tests  
      show-sql: true  
      properties:  
        hibernate:  
          format_sql: true  
          dialect: org.hibernate.dialect.PostgreSQLDialect
```

```

logging:
  level:
    org.hibernate.SQL: DEBUG
    org.testcontainers: INFO
    com.scriptmanager: DEBUG
  
```

### Key setting: `ddl-auto: create`

- `create` → Tables stay after test (can inspect with GUI)
  - `create-drop` → Tables dropped after test (clean slate)
- 

### 3. src/test/resources/junit-platform.properties

#### Content:

```

spring.test.constructor.autowire.mode=all
  
```

#### What it does:

- Enables constructor injection in tests
  - No need for `@Autowired` on constructor
  - Cleaner, more idiomatic Kotlin code
- 

## Test Configuration

### TestcontainersConfiguration.kt

**Location:** `src/test/kotlin/com/scriptmanager/config/TestcontainersConfiguration.kt`

#### Key features:

```

@TestConfiguration
class TestcontainersConfiguration {
    @Bean
    @ServiceConnection // ← Auto-configures datasource
    fun postgresContainer(): PostgreSQLContainer<*> {
        val container = PostgreSQLContainer(...)
            .withReuse(true) // ← Container reuse

        container.start()
        printConnectionInfo(container) // ← Shows connection details
        return container
    }
}
  
```

#### What it does:

1.  Starts PostgreSQL container automatically
  2.  Prints connection info (host, port, credentials)
  3.  Configures Spring datasource via `@ServiceConnection`
  4.  Reuses container for fast subsequent runs
- 

## DatabaseConfig.kt (Production Only)

**Important:** This config is **excluded from tests**:

```
@Configuration
@Profile("!test") // ← NOT active in test profile
class DatabaseConfig {
    @Bean
    fun dataSource(): DataSource {
        // SQLite for production
    }
}
```

**Why:** Tests use PostgreSQL from Testcontainers, not SQLite.

---

## Part 3: Writing Tests

### Basic Test Structure

#### Simple Test

```
@SpringBootTest
@Import(TestcontainersConfiguration::class)
@ActiveProfiles("test")
class CommandInvokerIntegrationTest(
    private val commandInvoker: CommandInvoker
) {
    @Test
    fun `context loads successfully`() {
        assertNotNull(commandInvoker)
    }
}
```

#### Key annotations:

- `@SpringBootTest` → Full Spring context
  - `@Import(TestcontainersConfiguration::class)` → PostgreSQL container
  - `@ActiveProfiles("test")` → Uses test configuration
  - Constructor parameters automatically injected (via `junit-platform.properties`)
-

## Test with Repository

```

@SpringBootTest
@Import(TestcontainersConfiguration::class)
@ActiveProfiles("test")
class WorkspaceIntegrationTest {
    private val commandInvoker: CommandInvoker,
    private val workspaceRepository: WorkspaceRepository
) {
    @Test
    fun `create workspace persists to database`() {
        // Execute
        commandInvoker.invoke(CreateWorkspaceCommand("Test"))

        // Verify
        val workspaces = workspaceRepository.findAll()
        assertEquals(1, workspaces.size)
        assertEquals("Test", workspaces[0].name.value)
    }
}

```

## Data Injection Tests

### Overview

Data injection tests allow you to set up test data once and reuse it across multiple test classes. This is useful for:

- Setting up reference data (e.g., model configs, users)
- Creating complex data relationships
- Avoiding duplicate setup code across tests
- Faster test execution (setup once, use many times)

### How Database Connection Sharing Works

Both **DataSetupTest** and other test classes share the **same PostgreSQL container** and **same database**:

```

@SpringBootTest
@Import(TestcontainersConfiguration::class) // ← Same Testcontainers
config
@ActiveProfiles("test") // ← Same profile

```

### Connection sharing mechanism:

- Container has **withReuse(true)** → Same container reused across tests
- Config has **ddl-auto: create** → Tables stay after first test

- Both use `@ServiceConnection` → Same database connection
- Data persists between test classes

**Flow:**

```

Test Suite Starts
  ↓
DataSetupTest runs
  ↓
Container starts (if not already running)
  ↓
Tables created (ddl-auto: create)
  ↓
Data injected to database
  ↓
DataSetupTest ends
  ↓
Tables STAY (ddl-auto: create, not create-drop)
  ↓
CommandInvokerIntegrationTest runs
  ↓
SAME container, SAME database
  ↓
Data from DataSetupTest is available 

```

**Key Configuration:**

```
# application-test.yml – MUST be 'create' not 'create-drop'
spring.jpa.hibernate.ddl-auto: create
```

**Why:**

- `create` → Tables stay after each test class (data persists)
- `create-drop` → Tables dropped after each test class (data lost!)

**DataSetupTest.kt**

**Purpose:** Inject test data that other tests can use.

**File Location:** `src/test/kotlin/com/scriptmanager/integration/DataSetupTest.kt`

**Key Features:**

- Runs first (controlled by test suite)
- Data persists to subsequent tests
- Controlled method execution order within class
- Shares instance across all methods

## Annotations Explained:

- `@TestMethodOrder(MethodOrderer.OrderAnnotation::class)` → Methods run in order (1, 2, 3...)
- `@TestInstance(TestInstance.Lifecycle.PER_CLASS)` → Share instance across all test methods
- `@Order(n)` → Specify execution order of methods

## Complete Example:

```

@SpringBootTest
@Import(TestcontainersConfiguration::class)
@ActiveProfiles("test")
@TestMethodOrder(MethodOrderer.OrderAnnotation::class) // ← Order methods
@TestInstance(TestInstance.Lifecycle.PER_CLASS) // ← Share
instance
class DataSetupTest(
    private val commandInvoker: CommandInvoker,
    private val workspaceRepository: WorkspaceRepository // Can inject
repositories too
) {
    private var workspaceId: Int? = null // Share data between methods

    @Test
    @Order(1)
    fun `01 - setup workspaces`() {
        println("🔧 Setting up test workspaces...")
        commandInvoker.invoke(CreateWorkspaceCommand("Test Workspace 1"))
        commandInvoker.invoke(CreateWorkspaceCommand("Test Workspace 2"))

        // Save ID for use in later methods
        val workspace = workspaceRepository.findAll().first()
        workspaceId = workspace.id

        println("✅ Workspaces created (ID: $workspaceId)")
    }

    @Test
    @Order(2)
    fun `02 - setup folders`() {
        println("🔧 Setting up test folders...")
        // Use workspaceId from previous method
        commandInvoker.invoke(
            CreateFolderCommand(
                name = "Test Folder",
                workspaceId = workspaceId!!
            )
        )
        println("✅ Folders created")
    }
}

```

`@Test`

```

@Order(3)
fun `03 - setup scripts`() {
    println("🔧 Setting up test scripts...")
    commandInvoker.invoke(
        CreateScriptCommand(
            name = "Test Script",
            content = "echo 'Hello from test'"
        )
    )
    println("✅ Scripts created")
}

@Test
@Order(4)
fun `04 - verify setup complete`() {
    println("✅ Verifying all test data...")
    val workspaces = workspaceRepository.findAll()
    assertTrue(workspaces.size >= 2, "Should have at least 2 workspaces")
    println("✅ Data setup verification complete!")
}
}

```

## Using Injected Data in Other Tests

After `DataSetupTest` runs, the data is available in all subsequent tests:

```

@SpringBootTest
@Import(TestcontainersConfiguration::class)
@ActiveProfiles("test")
class MyFeatureTest(
    private val workspaceRepository: WorkspaceRepository,
    private val scriptRepository: ScriptRepository
) {
    @Test
    fun `can query workspaces created by DataSetupTest`() {
        // Data from DataSetupTest is available!
        val workspaces = workspaceRepository.findAll()

        assertTrue(workspaces.size >= 2, "Should have workspaces from DataSetupTest")
        assertEquals("Test Workspace 1", workspaces[0].name.value)
    }

    @Test
    fun `can use existing scripts for testing new features`() {
        // Scripts created in DataSetupTest are available
        val scripts = scriptRepository.findAll()
        assertTrue(scripts.isNotEmpty(), "Should have scripts from DataSetupTest")
    }
}

```

```
// Test your new feature with existing data
    val result = myNewFeature.process(scripts.first())
    assertNotNull(result)
}
}
```

---

## Common Use Cases

### Use Case 1: Setup Reference Data

```
@Test
@Order(1)
fun `setup model configurations`() {
    commandInvoker.invoke(
        CreateModelConfigCommand(
            name = "gpt-4",
            apiKey = "test-key",
            endpoint = "https://api.openai.com"
        )
    )
    commandInvoker.invoke(
        CreateModelConfigCommand(
            name = "claude-3",
            apiKey = "test-key",
            endpoint = "https://api.anthropic.com"
        )
    )
}
```

### Use Case 2: Setup Test Users

```
@Test
@Order(1)
fun `setup test users`() {
    commandInvoker.invoke(
        CreateUserCommand(
            email = "admin@test.com",
            role = "ADMIN"
        )
    )
    commandInvoker.invoke(
        CreateUserCommand(
            email = "user@test.com",
            role = "USER"
        )
}
```

```
)  
}
```

### Use Case 3: Complex Data Relationships

```
private var workspaceId: Int? = null  
private var folderId: Int? = null  
  
@Test  
@Order(1)  
fun `01 - setup workspace`() {  
    commandInvoker.invoke(CreateWorkspaceCommand("Test Workspace"))  
    workspaceId = workspaceRepository.findAll().first().id  
}  
  
@Test  
@Order(2)  
fun `02 - setup folder in workspace`() {  
    // Uses workspaceId from previous test  
    commandInvoker.invoke(  
        CreateFolderCommand(  
            name = "Test Folder",  
            workspaceId = workspaceId!!  
        )  
    )  
    folderId = folderRepository.findAll().first().id  
}  
  
@Test  
@Order(3)  
fun `03 - setup scripts in folder`() {  
    // Uses folderId from previous test  
    commandInvoker.invoke(  
        CreateScriptCommand(  
            name = "Script 1",  
            folderId = folderId!!  
        )  
    )  
    commandInvoker.invoke(  
        CreateScriptCommand(  
            name = "Script 2",  
            folderId = folderId!!  
        )  
    )  
}
```

### Use Case 4: Bulk Data Creation

```

@Test
@Order(1)
fun `setup 100 test workspaces`() {
    repeat(100) { i ->
        commandInvoker.invoke(CreateWorkspaceCommand("Workspace $i"))
    }
    println("✅ Created 100 test workspaces for performance testing")
}

```

## Best Practices

### ✓ DO:

- ✅ Make data setup idempotent (can run multiple times safely)
- ✅ Use meaningful names for test data (e.g., "Test Workspace 1")
- ✅ Document what data is created in comments
- ✅ Use `@Order` annotations to control method execution
- ✅ Store IDs in class variables for use across methods
- ✅ Add verification tests to ensure data setup succeeded

### ✗ DON'T:

- ✗ Create too much data (slows down tests)
- ✗ Use hard-coded IDs (use returned IDs from commands)
- ✗ Assume order without `@Order` annotations
- ✗ Create data with sensitive or production-like values
- ✗ Forget to verify data was actually created

## Troubleshooting Data Injection

### Problem: Tests Run in Wrong Order

**Symptoms:** `DataSetupTest` runs after other tests

**Solution:** Always run via test suite:

```

./gradlew test --tests IntegrationTestSuite # ✅ Correct
./gradlew test                                # ✗ Order not guaranteed

```

### Problem: Data Not Available in Second Test

**Symptoms:** `CommandInvokerIntegrationTest` can't find data from `DataSetupTest`

**Possible causes:**

1. ✗ Using `create-drop` instead of `create` → Tables dropped between tests

2. ✗ Not running via test suite → Order not guaranteed
3. ✗ Transaction rollback → Data not committed
4. ✗ Container restarted → New database

## Solutions:

```
# Check application-test.yml
spring.jpa.hibernate.ddl-auto: create # ← MUST be 'create' not 'create-drop'
```

```
# Always run via test suite
./gradlew test --tests IntegrationTestSuite
```

## Problem: Container Not Shared

**Symptoms:** Each test starts a new container

**Solution:** Verify container reuse is enabled:

1. Check code:

```
// TestcontainersConfiguration.kt
.withReuse(true) // ← Must be set
```

2. Check global config:

```
cat ~/.testcontainers.properties
# Should contain: testcontainers.reuse.enable=true
```

3. Create if missing:

```
echo "testcontainers.reuse.enable=true" > ~/.testcontainers.properties
```

---

## Test Execution Order

### Overview

Controlling test execution order is crucial when you have tests that depend on data created by other tests. JUnit 5 provides several ways to control test order, and we use the **Test Suite** approach for maximum control.

## IntegrationTestSuite.kt

**Purpose:** Control which test classes run in which order.

**File Location:** `src/test/kotlin/com/scriptmanager/integration/IntegrationTestSuite.kt`

### Key Features:

- Explicit order control - tests run in the order specified
- Clear and maintainable - easy to see test execution flow
- Reliable - guaranteed execution order
- Flexible - easy to add or reorder tests

### Example:

```
@Suite
@SuiteDisplayName("Integration Tests with Data Setup")
@SelectClasses(
    DataSetupTest::class, // 1. Runs FIRST – injects
    data
    CommandInvokerIntegrationTest::class, // 2. Runs SECOND – uses data
    WorkspaceFeatureTest::class, // 3. Runs THIRD – more tests
    ScriptExecutionTest::class // 4. Runs FOURTH – even more
    tests
    // Add more test classes here in desired order
)
class IntegrationTestSuite
```

### Run the suite:

```
# Run all tests in correct order
./gradlew test --tests IntegrationTestSuite
```

### What happens:

1.  Container starts (or reuses existing)
2.  Schema created from JPA entities (if not exists)
3.  `DataSetupTest` runs → Injects data to database
4.  `CommandInvokerIntegrationTest` runs → Uses injected data
5.  `WorkspaceFeatureTest` runs → Uses same data
6.  `ScriptExecutionTest` runs → Uses same data
7.  All tests share same database connection and data

## Execution Order Options

### Option 1: Test Suite with `@SelectClasses` (Recommended )

**Use:** `@Suite` with `@SelectClasses`

**Example:**

```
@Suite
@SelectClasses(
    DataSetupTest::class,
    OtherTest::class
)
class IntegrationTestSuite
```

**Pros:**

- Explicit order control
- Clear which tests run in which order
- Reliable and deterministic
- Easy to maintain

**Cons:**

- Must maintain the list manually when adding new tests

**When to use:** When you need guaranteed execution order (recommended)

---

## Option 2: Class Order Annotation

**Use:** `@TestClassOrder(ClassOrderer.OrderAnnotation)` with `@Order` on classes

**Example:**

```
@Order(1)
class DataSetupTest { ... }

@Order(2)
class CommandInvokerIntegrationTest { ... }
```

**Pros:**

- Annotation-based ordering
- Works with test discovery

**Cons:**

- Requires JUnit 5.8+
- Less explicit than test suite
- Can be harder to see full test flow

**When to use:** When you prefer annotation-based configuration

### Option 3: Alphabetical Naming

**Use:** Name classes like `Test01_DataSetup`, `Test02_Integration`

**Example:**

```
class Test01_DataSetupTest { ... }
class Test02_CommandInvokerIntegrationTest { ... }
```

**Pros:**

- Simple naming convention

**Cons:**

- Fragile - depends on naming
- Not reliable across all test runners
- Makes test names awkward

**When to use:** Never (use Test Suite instead)

### Adding More Test Classes to Suite

To add a new test class to the execution order:

1. **Create your test class:**

```
@SpringBootTest
@Import(TestcontainersConfiguration::class)
@ActiveProfiles("test")
class MyNewFeatureTest(
    private val myService: MyService
) {
    @Test
    fun `test my new feature`() {
        // Test uses data from DataSetupTest
    }
}
```

2. **Add to IntegrationTestSuite:**

```
@Suite
@SelectClasses(
    DataSetupTest::class,           // 1. Data setup
    CommandInvokerIntegrationTest::class, // 2. Existing tests
    MyNewFeatureTest::class         // 3. Your new test ← Add here
```

```
)  
class IntegrationTestSuite
```

### 3. Run the suite:

```
./gradlew test --tests IntegrationTestSuite
```

---

## Running Tests Individually vs Suite

### Run Individual Test (Development)

```
# Run just data setup test  
./gradlew test --tests DataSetupTest  
  
# Run specific feature test  
./gradlew test --tests CommandInvokerIntegrationTest
```

#### Use when:

- Developing a specific test
- Quick feedback loop
- Debugging single test

**Note:** Data from DataSetupTest may not be available if you run other tests individually!

---

### Run Entire Suite (Recommended)

```
# Run all tests in correct order  
./gradlew test --tests IntegrationTestSuite
```

#### Use when:

- Running full test suite
  - CI/CD pipeline
  - Before committing code
  - Need guaranteed data availability
- 

### Run All Tests (No Order Guarantee)

```
# Runs all tests, but order NOT guaranteed  
./gradlew test
```

### Use when:

- $\Delta$  Tests are independent (don't rely on each other)
- $\Delta$  You understand the risks

**Warning:** Without test suite, execution order is not guaranteed! Tests may run in random order.

---

### Dependencies Between Tests

**Important:** JUnit 5 Test Suites don't have built-in dependency management. Tests in the suite run sequentially but:

- Each test is independent from JUnit's perspective
- Data sharing happens through database (not in-memory)
- If one test fails, subsequent tests still run

### Example flow:

```
DataSetupTest runs  
  ↓ (data saved to database)  
DataSetupTest passes/fails  
  ↓ (subsequent tests run regardless)  
CommandInvokerIntegrationTest runs  
  ↓ (reads data from database)  
CommandInvokerIntegrationTest passes/fails  
  ↓ (and so on....)
```

### If DataSetupTest fails:

- Data may not be created properly
  - $\Delta$  Subsequent tests may fail due to missing data
  - Check DataSetupTest first when debugging suite failures
- 

### Test Suite Best Practices

#### DO:

- Put data setup tests first in the suite
- Order tests from simple to complex
- Group related tests together
- Use descriptive suite display names
- Document test dependencies in comments

#### Example with documentation:

```
@Suite
@SuiteDisplayName("Complete Integration Test Suite")
@SelectClasses(
    // Phase 1: Data Setup
    DataSetupTest::class, // Creates workspaces, folders, scripts

    // Phase 2: Core Features
    CommandInvokerIntegrationTest::class, // Tests command invoker
    WorkspaceManagementTest::class, // Tests workspace CRUD

    // Phase 3: Advanced Features
    ScriptExecutionTest::class, // Tests script execution (needs scripts)
    AiIntegrationTest::class // Tests AI features (needs configs)
)
class IntegrationTestSuite
```

## ✗ DON'T:

- ✗ Mix test types in same suite (unit tests + integration tests)
- ✗ Create circular dependencies between tests
- ✗ Rely on test execution order for independent tests
- ✗ Add too many tests to one suite (split into multiple suites if needed)

## Multiple Test Suites

You can create multiple suites for different purposes:

```
// Fast tests – no data setup needed
@suite
@SelectClasses(
    SimpleUnitTest::class,
    QuickIntegrationTest::class
)
class FastTestSuite

// Full integration – with data setup
@suite
@SelectClasses(
    DataSetupTest::class,
    CompleteFeatureTest::class
)
class FullIntegrationTestSuite
```

Run specific suite:

```
./gradlew test --tests FastTestSuite  
./gradlew test --tests FullIntegrationTestSuite
```

---

## Part 4: Database Management

---

### Connecting to Test Database

#### Get Connection Info

When you run tests, connection info is printed:

```
=====  
=====  
🔗 TESTCONTAINERS DATABASE CONNECTION INFO  
=====  
=====  
📍 Host:      localhost  
⚡️ Port:      52106 ← USE THIS PORT!  
💾 Database:   testdb  
👤 Username:  test  
🔑 Password:  test  
🔗 JDBC URL: jdbc:postgresql://localhost:52106/testdb
```

Or use the script:

```
./get-db-connection.sh
```

---

### Connect with GUI Tools

#### DataGrip / IntelliJ Database

1. New Data Source → PostgreSQL
2. Host: **localhost** (NOT **http://localhost**)
3. Port: **52106** (from console output)
4. Database: **testdb**
5. User: **test**
6. Password: **test**

#### DBeaver / TablePlus / pgAdmin

Same settings as above.

**Common mistake:** Using **http://localhost** instead of just **localhost**

---

## Find Port Manually

```
# Check running containers
docker ps | grep postgres

# Output shows port mapping:
# 0.0.0.0:52106->5432/tcp
# The first number (52106) is your port
```

---

## Container Lifecycle

### How Container Reuse Works

#### First Test Run:

1. Container doesn't exist
2. Testcontainers creates PostgreSQL container
3. Assigns random port (e.g., 52106)
4. Container tagged for reuse
5. Test runs
6. Container stays alive 

#### Second Test Run:

1. Testcontainers looks for reusable container
2. Finds existing container
3. Reuses same container (same port!)
4. Test runs (much faster!)
5. Container stays alive 

**Key:** Container port stays the same until you manually stop it or restart Docker.

---

## Container Reuse Requirements

Both settings required:

### 1. In code:

```
.withReuse(true) // TestcontainersConfiguration.kt
```

### 2. Global config:

```
testcontainers.reuse.enable=true # ~/.testcontainers.properties
```

Without both, container is destroyed after each test run.

---

## Stopping Containers

### Option 1: Restart Docker (Simplest)

```
# Click Docker icon in menu bar → Restart  
# Or:  
osascript -e 'quit app "Docker"' && open -a Docker
```

#### Pros:

- Simplest method
- Cleans up everything
- No commands needed

#### Cons:

- ⚠ Stops ALL containers (not just test containers)
- 

### Option 2: Stop Test Container Only

```
# Use the script  
./stop-test-db.sh  
  
# Or manually  
docker stop $(docker ps -q --filter ancestor=postgres:15-alpine)  
docker rm $(docker ps -aq --filter ancestor=postgres:15-alpine)
```

#### Pros:

- Only stops test container
  - Other containers keep running
- 

### Option 3: Automatic Stop (Change Config)

In `TestcontainersConfiguration.kt`:

```
.withReuse(false) // Container stops after tests
```

#### Pros:

- Fully automatic
- No manual cleanup

**Cons:**

- Slower test runs (container recreated each time)
- 

**When to Stop****Stop when:**

- Done testing for the day
- Need to free the port
- Want clean slate

**Keep running when:**

- Still actively testing
  - Want fast test runs
  - Inspecting database with GUI
- 

## Part 5: Understanding Your Setup

---

### Spring vs Prisma

Question: Is there Spring + Prisma integration?

**Answer: NO **

Technology	Ecosystem	ORM
<b>Prisma</b>	Node.js/TypeScript	Prisma ORM
<b>Spring</b>	Java/Kotlin	JPA/Hibernate

**Your setup:**

- Spring Boot with JPA/Hibernate
  - 16 JPA entities with `@Entity` annotations
  - Hibernate generates schema from entities
  - Prisma NOT used in tests (only in production if needed)
- 

### Why Not Use Prisma for Tests?

**Problems:**

- Requires Node.js/npm
- Prisma schema must be manually synced with JPA entities
- Two sources of truth (Prisma schema + JPA entities)
- Extra complexity

## Solution (Current):

- JPA entities are single source of truth
  - Hibernate auto-generates schema from entities
  - No external dependencies
  - Schema always matches entities
- 

## JPA Foreign Keys

Question: Can JPA automatically create foreign keys like Prisma?

Answer: YES!

### How it works:

```
// Your JPA entity
@Entity
class ScriptsFolder {
    @OneToMany
    @JoinTable(
        name = "rel_workspace_folder",
        joinColumns = [JoinColumn(name = "workspace_id")],
        inverseJoinColumns = [JoinColumn(name = "folder_id")]
    )
    var folders: MutableSet<ScriptsFolder>
}
```

### Hibernate generates:

```
CREATE TABLE rel_workspace_folder
(
    workspace_id INTEGER NOT NULL,
    folder_id    INTEGER NOT NULL,
    CONSTRAINT fk_workspace
        FOREIGN KEY (workspace_id) REFERENCES workspace (id),
    CONSTRAINT fk_folder
        FOREIGN KEY (folder_id) REFERENCES scripts_folder (id)
);
```

All foreign key constraints are automatically created!

---

## JPA Relationship Annotations

Annotation	Creates	Use For
@ManyToOne	FK column	Many-to-one relationship

Annotation	Creates	Use For
@OneToMany	FK in other table	One-to-many relationship
@JoinColumn	Direct FK	Single foreign key
@JoinTable	Join table + 2 FKs	Many-to-many relationship

### Example from your entities:

- Workspace → Folders (join table)
- Folder → Scripts (join table)
- Folder → Subfolder (self-referencing)
- All cascade operations work

## SQLite vs PostgreSQL in Tests

Question: Why SQLite error in tests?

**Problem:** Test was using SQLite instead of PostgreSQL.

**Root cause:** `DatabaseConfig` was active in test profile, creating SQLite datasource.

**Solution:** Added `@Profile("!test")` to exclude it from tests:

```
@Configuration
@Profile("!test") // ← Only active when NOT test
class DatabaseConfig {
    @Bean
    fun dataSource(): DataSource {
        // SQLite for production
    }
}
```

**Result:**

- Production: Uses SQLite
- Tests: Uses PostgreSQL (from Testcontainers)

## Why Different Databases?

**Benefits:**

1.  PostgreSQL has full SQL support (better for tests)
2.  SQLite is lightweight (good for desktop app)
3.  Tests catch SQL compatibility issues
4.  JPA abstracts away database differences

**Your setup:**

- Production: SQLite (configured in DatabaseConfig)
  - Tests: PostgreSQL (configured by Testcontainers)
- 

## Tables Dropping After Tests

Question: Why do tables drop after tests?

**Answer:** This is expected behavior based on `ddl-auto` setting.

### Options

```
# Option 1: create-drop (Clean slate)
spring.jpa.hibernate.ddl-auto: create-drop
```

- Creates tables on start
- Drops tables on end
- Clean every run

```
# Option 2: create (Keep for inspection) – CURRENT
spring.jpa.hibernate.ddl-auto: create
```

- Creates tables on start
- Keeps tables on end
- Can inspect with GUI

**Current setting:** `create` (tables stay after tests)

---

### What Happens

**With `create-drop`:**

```
Test starts → Creates tables → Tests run → Drops tables → Clean
```

**With `create (current)`:**

```
Test starts → Creates tables → Tests run → Keeps tables → Can inspect
```

**On next run:**

```
Drops old tables → Creates new tables → Fresh data
```

## Part 6: Troubleshooting

---

### Common Issues

#### Issue 1: "Docker not running"

**Symptoms:**

```
Error: Could not find Docker environment
```

**Solution:**

```
open -a Docker # Start Docker  
docker ps      # Verify
```

---

#### Issue 2: "Container not reused"

**Symptoms:** New container created every test run (slow)**Solution:**

1. Check `~/.testcontainers.properties` exists:

```
cat ~/.testcontainers.properties  
# Should show: testcontainers.reuse.enable=true
```

2. Create if missing:

```
echo "testcontainers.reuse.enable=true" > ~/.testcontainers.properties
```

---

#### Issue 3: "Constructor injection not working"

**Symptoms:**

```
No ParameterResolver registered for parameter [...]
```

**Solution:** Check `src/test/resources/junit-platform.properties` exists:

```
spring.test.constructor.autowire.mode=all
```

---

## Issue 4: "SQLite error in tests"

**Symptoms:**

```
org.sqlite.SQLiteException: [SQLITE_ERROR] SQL error
```

**Solution:** Verify DatabaseConfig has `@Profile("!test")`:

```
@Configuration  
{@Profile("!test") // ← Must be present  
class DatabaseConfig { ... }}
```

---

## Issue 5: "Data not available in second test"

**Symptoms:** CommandInvokerIntegrationTest can't find data from DataSetupTest**Possible causes:**

1. ❌ Using `create-drop` instead of `create`
2. ❌ Not running via test suite
3. ❌ Wrong execution order

**Solution:**

1. Check `application-test.yml`:

```
spring.jpa.hibernate.ddl-auto: create # NOT create-drop
```

2. Run via test suite:

```
./gradlew test --tests IntegrationTestSuite
```

---

## Issue 6: "Can't connect with GUI tool"

**Symptoms:** Connection refused**Common mistakes:**

- ✗ Using `http://localhost` → Should be just `localhost`
- ✗ Using port 5432 → Should be mapped port (e.g., 52106)
- ✗ Container stopped → Check with `docker ps | grep postgres`

**Solution:**

1. Get correct port:

```
./get-db-connection.sh
```

2. Use settings from output:

- Host: `localhost` (no `http://`)
  - Port: from script output
  - Database: `testdb`
  - User: `test`
  - Password: `test`
- 

## FAQ

Q: How do I run tests?

```
./gradlew test
```

Q: How do I run tests with data injection?

```
./gradlew test --tests IntegrationTestSuite
```

Q: How do I connect to the test database?

```
./get-db-connection.sh
```

Then use the connection info in your GUI tool.

---

Q: How do I stop the test container?

**Option 1 (Simplest):** Restart Docker **Option 2:** Run `./stop-test-db.sh` **Option 3:** Change config to `withReuse(false)`

---

Q: Do I need Docker for tests?

**Yes**, Docker must be running for integration tests with Testcontainers.

---

Q: Can I use SQLite for tests instead?

**Not recommended.** SQLite has limited SQL features. PostgreSQL provides:

- Full SQL support
  - Foreign key constraints
  - Production-like behavior
- 

Q: Why does the port change?

The port is **dynamically assigned** by Docker. But with `withReuse(true)`, the same container (and same port) is reused until you stop it.

---

Q: How do I add more test data?

Edit `DataSetupTest.kt` and add your commands:

```
@Test  
@Order(1)  
fun `01 – setup workspaces`() {  
    commandInvoker.invoke(CreateWorkspaceCommand("My Workspace"))  
}
```

---

Q: How do I verify foreign keys were created?

1. Run test
2. Connect with GUI tool (DataGrip, DBeaver)
3. View table structure
4. Check "Foreign Keys" tab

Or run SQL:

```
SELECT *  
FROM information_schema.table_constraints  
WHERE constraint_type = 'FOREIGN KEY';
```

---

Q: What's the difference between `create` and `create-drop`?

Setting	Tables After Test	Use For
---------	-------------------	---------

Setting	Tables After Test	Use For
create	Stay	Development (can inspect)
create-drop	Dropped	CI/CD (clean slate)

Q: Why 16 entities? What are they?

You have 16 JPA entities in your project:

- Workspace, ScriptsFolder, ShellScript
- AiProfile, ModelConfig, ScriptAiConfig
- Event, HistoricalShellScript
- And 8 more...

All automatically get tables + foreign keys created by Hibernate!

## Quick Commands Reference

```
# Start Docker
open -a Docker

# Run all tests
./gradlew test

# Run test suite with data injection
./gradlew test --tests IntegrationTestSuite

# Run specific test
./gradlew test --tests CommandInvokerIntegrationTest

# Get database connection info
./get-db-connection.sh

# Stop test container
./stop-test-db.sh

# Create testcontainers config
./create-testcontainers-config.sh

# Check container is running
docker ps | grep postgres

# Restart Docker (stops all containers)
# Click Docker icon → Restart
```

---

## File Reference

## Configuration Files

- `src/test/resources/application-test.yml` - Test config
- `src/test/resources/junit-platform.properties` - Constructor injection
- `~/.testcontainers.properties` - Container reuse

## Test Files

- `src/test/kotlin/com/scriptmanager/integration/`
  - `DataSetupTest.kt` - Data injection
  - `CommandInvokerIntegrationTest.kt` - Example test
  - `IntegrationTestSuite.kt` - Test execution order

## Config Classes

- `src/test/kotlin/com/scriptmanager/config/`
  - `TestcontainersConfiguration.kt` - Container setup

## Helper Scripts

- `get-db-connection.sh` - Get connection info
- `stop-test-db.sh` - Stop container
- `create-testcontainers-config.sh` - Setup config

---

## Summary

### What You Have

✓ Spring Boot 3.2 with Kotlin ✓ PostgreSQL via Testcontainers ✓ 16 JPA entities with automatic schema generation ✓ Container reuse for fast tests ✓ Data injection test pattern ✓ Test execution order control ✓ Foreign keys automatically created ✓ GUI connection support

### Key Features

- 🚀 **Fast:** Container reuse makes subsequent runs seconds
- 🔒 **Production-like:** Real PostgreSQL, not in-memory
- ✎ **Isolated:** Clean database for each test suite run
- 📊 **Inspectable:** Connect with GUI tools
- ⏱ **Controlled:** Test execution order guaranteed

### Next Steps

1. **Run tests:** `./gradlew test`
2. **Add test data:** Edit `DataSetupTest.kt`
3. **Inspect database:** Use `./get-db-connection.sh` + GUI tool
4. **Write more tests:** Follow the patterns in the guide

---

🎉 You're all set! Everything is configured and ready to use!

**Questions?** Refer to the [Troubleshooting](#) and [FAQ](#) sections.

---

*This guide consolidates all testing and Testcontainers documentation. No other testing guides needed.*